

Satisfiability modulo theories Part 2

Neural Theory Solvers

Verifying cyberphysical systems

Sayan Mitra

mitras@illinois.edu

Today

- SMT
- Decision procedure for Linear Real Arithmetic
Simplex Algorithm [Dantzig 1947]
- Next week: Verification of Neural Networks
Reluplex [Katz et al 2017]

References

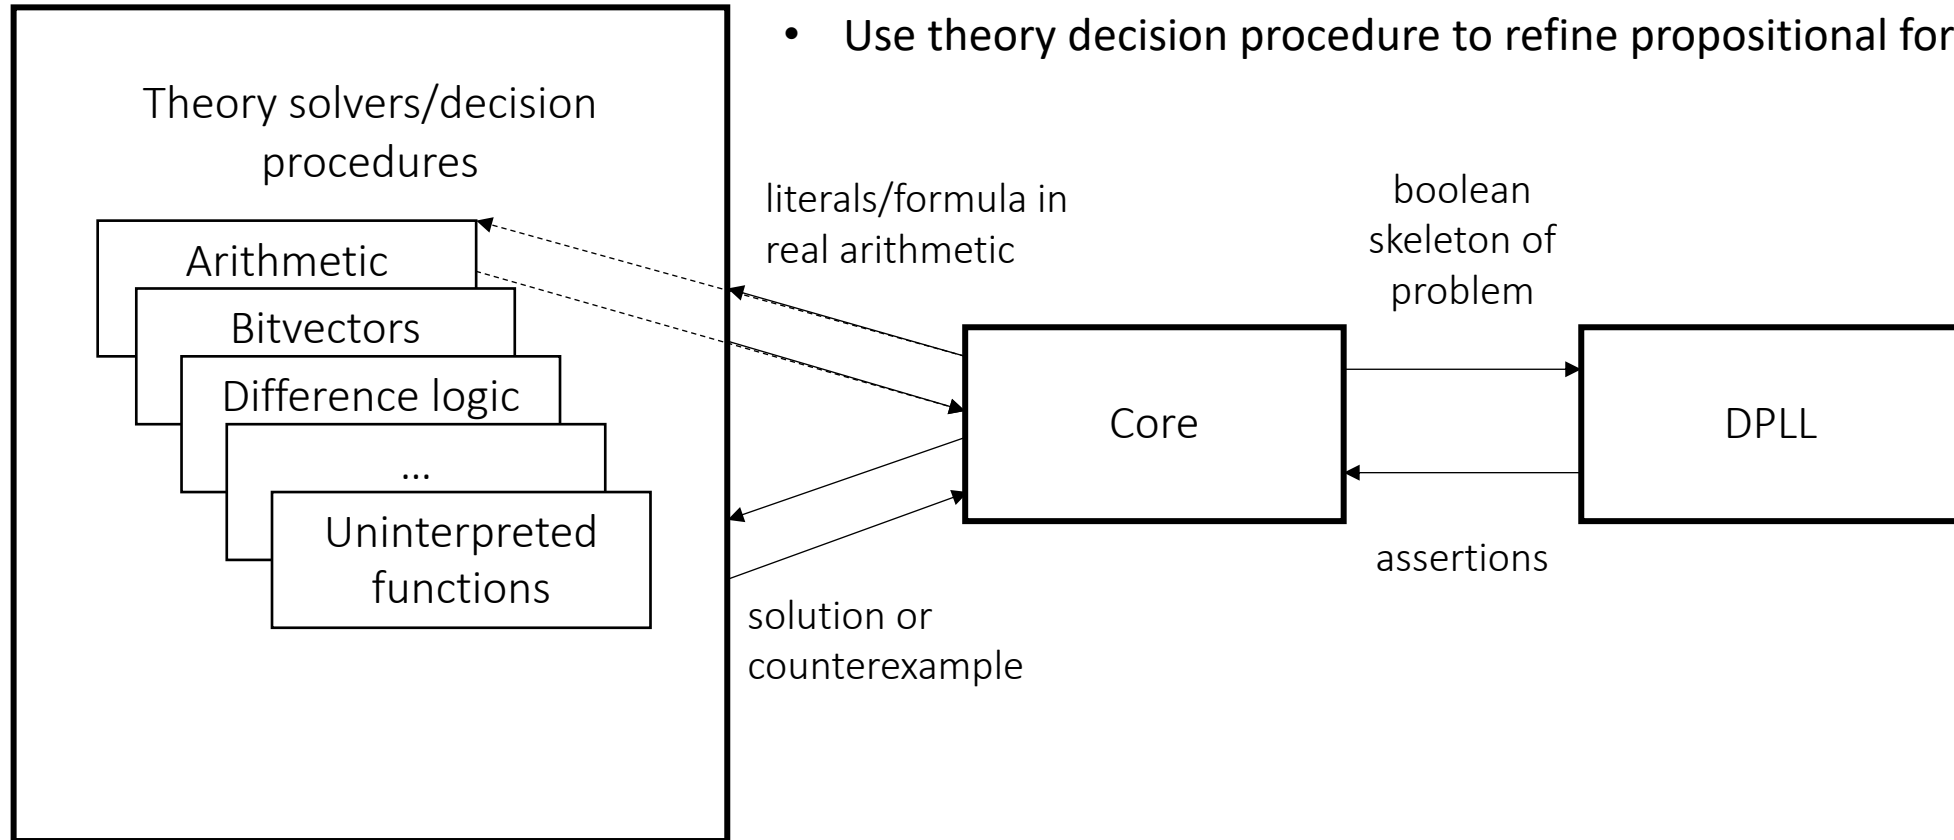
- Lectures on SMT from Clark Barrett
- Book: Introduction to Neural Network Verification by Aws Albarghouthi
- Book: Decision Procedures by Daniel Kroening and Ofer Strichman

SMT

$$\phi \equiv g(a) = c \wedge f(g(a)) \neq f(c) \vee g(a) = d \wedge c \neq d$$

Several approaches, lazy approach:

- Abstract ϕ to propositional form
- Feed to DPLL
- Use theory decision procedure to refine propositional formula a guide SAT



DPLL^T: DPLL modulo theories

How can we extend DPLL to handle formulas over other theories like

- Difference Logic (DL)
- Linear Real Arithmetic (LRA)
- Uninterpreted functions (UF)

Idea: Start with a *Boolean abstraction* (or skeleton) and incrementally add more *theory* information until we can conclusively say SAT or UNSAT

Example: DPLL^{LRA}

$$F \equiv (x \leq 0 \vee x \leq 10) \wedge (\neg x \leq 0)$$

Boolean abstraction: replace every unique linear inequality with a Boolean variable

$$F^B \equiv (p \vee q) \wedge (\neg p)$$

where p abstracts $x \leq 0$ and q abstracts $x \leq 10$

Abstraction because information is lost

The relationship $x > 10 \Rightarrow x > 0$, i.e., $\neg q \Rightarrow \neg p$ is lost in F_B

Notation. $(F^B)^T$ maps F^B back to theory T , i.e., $(F^B)^T = F$.

Proposition. If F^B is UNSAT then F is UNSAT, but the converse does not hold, i.e., F^B is SAT does not mean that F is SAT.

Example. $F_1 \equiv (x \leq 0 \wedge x \geq 10)$ is clearly UNSAT, however $F_1^B \equiv p \wedge q$ is SAT.

Lazy DPLL^T Algorithm using a Decision Procedure $T()$

Input: A formula F in CNF form over theory T

Output: $I \models F$ or UNSAT

Let F^B be the abstraction of F

while true do

if DPLL(F^B) is unsat **then return** UNSAT

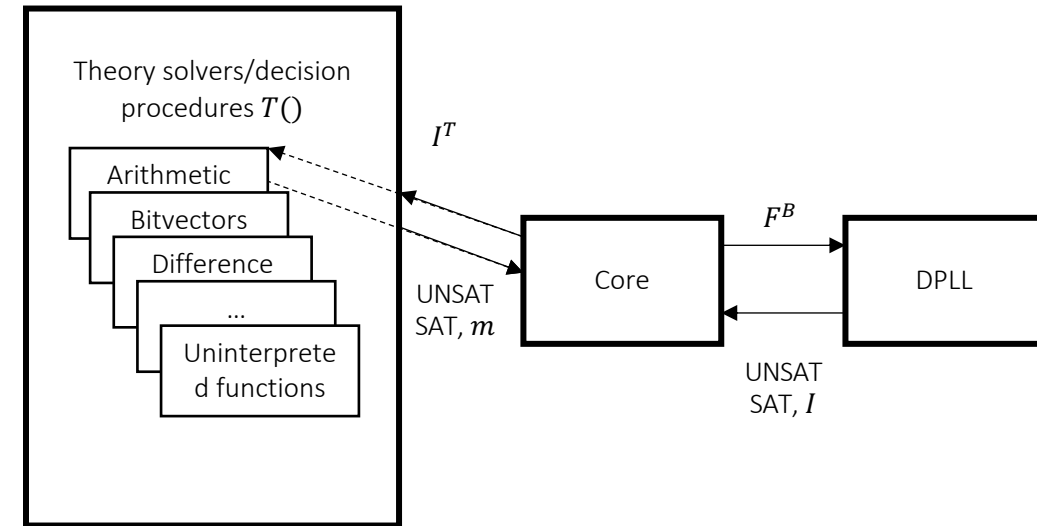
else

 Let I be the model returned by *DPLL*

 Assume I is represented as a formula

if $T(I^T)$ is sat **then return** SAT and the model returned by $T()$

else $F^B := F^B \wedge \neg I$



- $\phi \equiv \underbrace{g(a) = c}_1 \wedge \underbrace{f(g(a)) \neq f(c)}_{\bar{2}} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\bar{4}}$

- send $\phi^B \equiv \{1, \bar{2} \vee 3, \bar{4}\}$ to DPLL

- DPLL returns SAT with model $I: \{1, \bar{2}, \bar{4}\}$

- UF solver concretizes $I^{UF} \equiv g(a) = c, f(g(a)) \neq f(c), c \neq d$

- UF checks I^{UF} as UNSAT

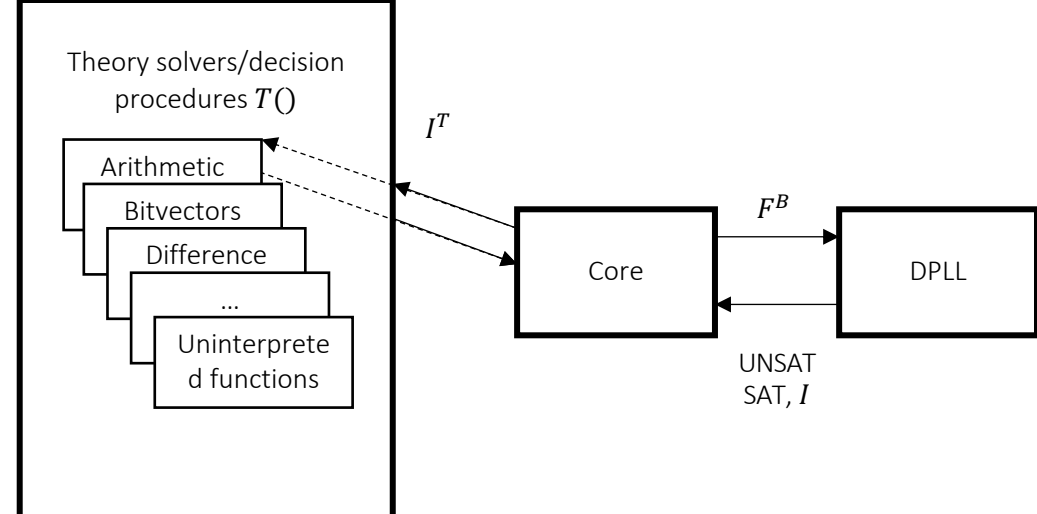
- send $\phi^B \wedge \neg I: \{1, \bar{2} \vee 3, \bar{4}, \bar{1} \vee 2 \vee 4\}$ to DPLL; this is a new fact learned by DPLL

- DPLL returns model $I': \{1, 2, 3, \bar{4}\}$

- UF solver concretizes I'^{UF} and finds this to be UNSAT

- send $\phi^B \wedge \neg I \wedge \neg I': \{1, \bar{2} \vee 3, \bar{4}, \bar{1} \vee 2 \vee 4, \bar{1} \vee \bar{2} \vee \bar{3} \vee 4\}$ to DPLL; another fact

- returns UNSAT



Neural Theory Solvers Simplex and ReluPlex

Verifying cyberphysical systems

Sayan Mitra

mitras@illinois.edu

Reference : Introduction to Neural Network Verification by Aws Albarghouthi

Decision Procedure for Linear Real Arithmetic

Input: $F \equiv \bigwedge_{i=1}^n \sum_{j=1}^m c_{ij} x_j \leq b_i$ where $c_{ij}, b_i \in \mathbb{R}$

Output: $\exists \mathbf{x} \in \mathbb{R}^m$ such that $\mathbf{x} \models F$?

Solution based on Simplex Algorithm [Dantzig 1947]

Simplex solves

$\max \sum_{j=1}^m a_j x_j$ subject to

$$\bigwedge_{i=1}^n \sum_{j=1}^m c_{ij} x_j \leq b_i$$

Our focus will be on finding any solution $\mathbf{x} \in \mathbb{R}^m$ that satisfies F

Decision Procedure for Linear Real Arithmetic

Input: $F \equiv \bigwedge_{i=1}^n \sum_{j=1}^m c_{ij} x_j \leq b_i$ where $c_{ij}, b_i \in \mathbb{R}$

Output: \exists a model $\mathbf{x} \in \mathbb{R}^m$ such that $\mathbf{x} \models F$?

Simplex expects F to be expressed in the Simplex form, which is a conjunction of

- Linear equalities: $\sum_{i=1}^m c_i x_i = 0$
- Bounds: $l_i \leq x_i \leq u_i$

Transforming to Simplex Form

Consider the i^{th} inequality in F : $\sum_{j=1}^m c_{ij}x_j \leq b_i$

Rewrite this as:

$$s_i = \sum_{j=1}^m c_{ij}x_j \wedge$$

$$s_i \leq b_i$$

s_i is called a *slack variable*

Putting together all the rewritten conjuncts we get F_S

Proposition.


1. Any model of F_S is a model of F , disregarding the assignments to the slack variables.
2. If F_S is UNSAT then F is UNSAT.

Simplex (Informal)

Idea. Simultaneously try to find a model or a proof of UNSAT

Start with some *model (or valuation)* that satisfies all linear equalities (say, $x_i = 0, \forall i$)

In each iteration, pick a bound that is not satisfied and modify the model to satisfy the bound OR discover that the formula is UNSAT



$$x_0 = \langle x \mapsto 0, y \mapsto 0 \rangle$$

$$x_0 \setminus \text{unsat } -2x + y \geq 2$$

$$x_1 = \langle x \mapsto -1, y \mapsto 0 \rangle$$

$$x_1 \setminus \text{unsat } x + y \geq 0$$

$$x_2 = \left\langle x \mapsto -\frac{1}{2}, y \mapsto 1 \right\rangle \models F$$

Variable naming and ordering for Simplex

The input formula F_S (after rewriting) has two types of variables

- **Basic variables** appear on the LHS of an equality; initially these are the *slack variables*
- **Non-basic variables** all others

In each iteration, some basic variable becomes non-basic

We fix an *arbitrary total ordering* on variables x_1, \dots, x_n

For a basic variable x_i and non-basic variable x_j we denote by c_{ij} the coefficient of x_j in the definition of x_i , i.e.,

$$x_i = \dots + c_{ij} x_j + \dots$$

The upper and lower bounds of x_i are called u_i and l_i (possibly $\infty, -\infty$)

Simplex (Formal) 1

The algorithm maintains two invariants

1. The model x always satisfies the equalities; bounds may be violated.

Why is this invariant satisfied by our initialization of all 0s?

2. The bounds of all non-basic variables are all satisfied.

Why is this invariant satisfied by our initialization?

Simplex Algorithm: DP for LRA

Input: A formula F_S in Simplex form

Output: $x \models F_S$ or UNSAT

$x := \langle x_i \mapsto 0 \rangle$

while true **do**

if $x \models F_S$ **then return** x

Let x_i be the first basic variable s.t. $x[x_i < l_i$ or $x[x_i > u_i$

if $x[x_i < l_i$ **then**

Let x_j be the first non-basic variable s.t.

$(x[x_j < u_j \wedge c_{ij} > 0]) \vee (x[x_j > l_j \wedge c_{ij} < 0])$

If no such x_j exists **then return** UNSAT

$x[x_j := x[x_j + \frac{l_i - x[x_i]}{c_{ij}}$

else Let x_j be the first non-basic variable s.t.

$(x[x_j > l_j \wedge c_{ij} > 0]) \vee (x[x_j < u_j \wedge c_{ij} < 0])$

If no such x_j exists **then return** UNSAT

$x[x_j := x[x_j + \frac{u_i - x[x_i]}{c_{ij}}$

Pivot x_i and x_j

$$x_i = \sum_{k \in N}^m c_{ik} x_k, j \in N$$

Pivoting x_i and x_j rewrites x_j as basic variable

$$x_i = c_{ij} x_j + \sum_{k \in N \setminus \{j\}}^m c_{ik} x_k$$

$$x_j = -\frac{x_i}{c_{ij}} + \sum_{k \in N \setminus \{j\}}^m \frac{c_{ik}}{c_{ij}} x_k$$

Example

$$x + y \geq 0$$

$$-2x + y \geq 2$$

$$-10x + y \geq -5$$

Rewritten in Simplex form

$$s_1 = x + y$$

$$s_2 = -2x + y$$

$$s_3 = -10x + y$$

$$s_1 \geq 0$$

$$s_2 \geq 2$$

$$s_3 \geq -5$$

Example continued

Variable ordering

x, y, s_1, s_2, s_3

Initialization $\mathbf{x}_0 = \langle x \mapsto 0, y \mapsto 0, s_1 \mapsto 0, s_2 \mapsto 0, s_3 \mapsto 0 \rangle$

\mathbf{x}_0 satisfies equalities, bounds of s_1 s_3 are satisfied

Pick the first variable x to fix the bound of s_2

Since upper and lower bounds of x are ∞ and $-\infty$ it easily satisfies the blue condition

To increase s_2 to 2 and satisfy its lowerbound we decrease x [x to -1

$\mathbf{x}_1 = \langle x \mapsto -1, y \mapsto 0, s_1 \mapsto -1, s_2 \mapsto 2, s_3 \mapsto 10 \rangle$

Pivot s_2 with x

$$s_1 = x + y$$

$$s_2 = -2x + y$$

$$s_3 = -10x + y$$

$$s_1 \geq 0$$

$$s_2 \geq 2$$

$$s_3 \geq -5$$

$$x = -0.5s_2 + 0.5y$$

$$s_1 = -0.5s_2 + 1.5y$$

$$s_3 = 5s_2 - 4y$$

$$s_1 \geq 0$$

$$s_3 \geq -5$$

$$-\infty \leq x \leq \infty$$

Example continued 2

$$\begin{aligned}x &= -0.5s_2 + 0.5y \\s_1 &= -0.5s_2 + 1.5y \\s_3 &= 5s_2 - 4y \\s_1 &\geq 0 \\s_3 &\geq -5 \\-\infty &\leq x \leq \infty\end{aligned}$$

$$x_1 = \langle x \mapsto -1, y \mapsto 0, s_1 \mapsto -1, s_2 \mapsto 2, s_3 \mapsto 10 \rangle$$

All equalities are still satisfied (invariant)

The only basic variable not satisfying its bounds is now s_1

The first non-basic variable we can tweak is y

Setting $y=1$ to satisfy the lowerbound of s_1 we get

$$x_2 = \langle x \mapsto -0.5, y \mapsto 1, s_1 \mapsto 0.5, s_2 \mapsto 2, s_3 \mapsto 6 \rangle$$

Pivot s_1 with y

$$x_2 \models F_S$$

$$\begin{aligned}y &= \frac{2}{3}s_1 + \frac{1}{3}s_2 \\x &= +\frac{1}{3}s_1 - \frac{1}{3}s_2 \\s_2 &\geq 2 \\s_1 &\geq 0 \\s_3 &\geq -5 \\-\infty &\leq x \leq \infty\end{aligned}$$

Why is simplex correct?

- Why does it terminate?

Because we always look for the first variable violating the bounds. There is a property (Bland's rule) that ensures that we never revisit the same set of basic and non-basic variables.

- Why does it give the right answer (sound)?

- If it returns x does it satisfy $x \models F$?

This follows from the condition before **return** x

- If it returns UNSAT is F really unsatisfiable?

Unsatisfiable example

$$s_1 = x + y$$

$$s_2 = -x - 2y$$

$$s_3 = -x + y$$

$$s_1 \geq 0$$

$$s_2 \geq 2$$

$$s_3 \geq 1$$

Consider a Simplex execution in which there are two pivots:

Pivot 1: s_1 with x

$$x = s_1 - y$$

$$s_2 = -s_1 - y$$

$$s_3 = -s_1 + 2y$$

Pivot 2: s_2 with y

$$x = 2s_1 + s_2$$

$$y = -s_1 - s_2$$

$$s_3 = -3s_1 - 2s_2$$

Non-basic variables satisfy their bounds (invariant) and so $s_1 \geq 0, s_2 \geq 2$

If s_2 violates the bound then

$$s_3 = -3s_1 - 2s_2 < 1$$

We can make s_3 bigger by decreasing s_1 and s_2 but the at most

$$s_3 = -3.0 - 2.2 = -4$$

which is still less than 1 and Simplex concludes that the formula is UNSAT.

The blue conditions for choosing x_j encodes this condition.

Assignments

- Learn z3
 - <https://ericpony.github.io/z3py-tutorial/guide-examples.htm>

Readings

- Read chapter 4 for next week
- Reading more about decision procedures

